

LovelaceJEPA: Online Encoder Adaptation by Learning Anchored Consistent Embeddings

Bourke Floyd IV
Independent Researcher
bourke.floyd@example.com*

Abstract

Test-time adaptation (TTA) lets a latent world model track distribution shift inside the closed loop of model-predictive control: after each executed action, the observed transition becomes a self-supervised target and the model takes one gradient step. But the standard objective is self-referential — the target $\text{sg}(E_\theta(o_{t+1}))$ is produced by the very encoder being adapted — so adaptation can reduce prediction error by *relocating the latent space* rather than by modeling the new dynamics. Every consumer calibrated on the pretrained space — goal embeddings, success and progress heads, the planner’s cost — is silently damaged, while the adaptation loss looks excellent. We observed this failure on a deployed screen-interaction agent: adaptation cut prediction divergence by 45% in the second half of each run while the frozen success head’s AUC fell from 0.933 to 0.826. We propose LACE (Learning Anchored Consistent Embeddings): replace the adaptation target with $E_{\bar{\theta}}(o_{t+1})$ from a frozen (or slow-EMA) copy of the pretrained encoder, anchoring adaptation to the pretrained manifold. The change is one symbol in the objective and adds no planning-time cost. On miniature reproductions of the PointMaze shift suite and a pymunk pusher benchmark, LACE keeps the planning gains of self-referential TTA while eliminating goal-latent drift entirely and surviving symmetric learning rates that collapse the unlaced objective. On 110 replayed deployment runs (5,434 steps), anchoring the encoder is necessary but not sufficient: prequential replay decomposes the head damage into an encoder-relocation component that LACE removes and a predictor-drift component that persists, which we isolate with parameter-subset ablations.

1 Introduction

“The Analytical Engine has no pretensions whatever to originate anything. It can do whatever we know how to order it to perform.” — Ada Lovelace, Note G (1843)

Latent world models paired with model-predictive control (MPC) act by imagining futures: an encoder E maps observations to latents, a predictor f rolls latents forward under actions, and a planner selects the action sequence whose imagined future best matches a goal embedding [9, 10, 20]. When the deployed environment drifts from the pretraining distribution, recent work adapts the world model *online*, inside the control loop: AdaJEPA [1] takes one gradient step per replan on the self-supervised latent prediction loss

$$\mathcal{L}_{\text{ada}} = \|f(E_\theta(o_t), a_t) - \text{sg}(E_\theta(o_{t+1}))\|^2, \quad (1)$$

updating a small parameter subset, and recovers a large fraction of the planning success lost to dynamics, visual, and geometry shifts.

This paper starts from a deployment observation that complicates the story. Equation (1) is *self-referential*: the target is produced by the very encoder being adapted. Nothing in the objective prefers “model the new dynamics” over “move the latent space somewhere easier to predict.” The loss cannot see the difference — but every frozen consumer of the latent space can. Goal latents drift although the goal observation never changes; frozen success/progress heads receive off-manifold inputs; the planning cost is evaluated in a space that no longer means what it meant at calibration time. Running the published recipe

*Placeholder contact; update before submission.

on a deployed screen-interaction agent (Section 5.8), adaptation reduced second-half prediction divergence by 45% — and dropped the deployed success detector’s AUC from 0.933 to 0.826. The standard defense is to adapt the encoder much more slowly than the predictor (AdaJEPA uses $\eta_{\text{enc}} = 10^{-5}$ vs $\eta_{\text{pred}} = 5 \times 10^{-4}$), which limits but does not remove the damage, and sacrifices adaptation capacity exactly where visual shifts need it.

Our fix is deliberately minimal. LACE (Learning Anchored Consistent Embeddings) changes one symbol in Equation (1): the target comes from a *frozen* copy of the pretrained encoder,

$$\mathcal{L}_{\text{LACE}} = \|f(E_{\theta}(o_t), a_t) - E_{\bar{\theta}}(o_{t+1})\|^2, \quad (2)$$

(or a slow EMA of the student). The student encoder and predictor may now adapt at full strength: any relocation of the latent space is penalized because the target stays put. A corollary falls out for free: the goal should be encoded by the anchor, pinning the planner’s cost landscape to the space its consumers were calibrated on. We call the resulting model LovelaceJEPA, because its adaptation “has no pretensions to originate anything”: it may only re-express the pretrained representation, never invent a new one.

Contributions. (1) We identify and name the *head-consistency failure* of self-referential TTA in latent world models, and make it measurable on public benchmarks via frozen probe heads (Section 4.1). (2) We propose LACE, a one-line anchored-target objective with a goal-encoding corollary (Section 4). (3) On miniature reproductions of AdaJEPA’s benchmarks (PointMaze shift suite; a pymunk pusher suite) we show that anchoring costs no planning success, eliminates goal-latent drift, and survives aggressive symmetric learning rates that collapse the self-referential objective (Section 5). (4) On 110 replayed runs (5,434 steps) of a deployed screen-interaction agent, we decompose head damage into an encoder-relocation component that anchoring removes and a predictor-drift component that it cannot, isolating each with parameter-subset ablations (Section 5.8).

2 Related work

Latent world models and MPC. World models learned by joint-embedding predictive architectures (JEPA) [13, 2] plan in latent space without reconstruction [11, 20]. Our setting and baseline follow AdaJEPA [1], which adapts a JEPA world model online during MPC.

Test-time training and adaptation. TTT with self-supervised auxiliary tasks [16, 7] and entropy minimization [17] adapt perception models at test time; the collapse risk of self-referential objectives is managed there with restricted parameter subsets and small steps. The continual-TTA literature has developed source-anchoring devices against long-horizon drift — weight-averaged teachers with stochastic restore [18], Fisher-weighted regularization toward source weights [14], self-distillation from a frozen copy of the source network [15, 5] — and Zhao et al. [19] document that TTA is bounded by the frozen classifier it feeds. All of this work anchors *classifier logits or weights* of a perception model under an entropy or pseudo-label objective. Our setting differs in both loop and objective: the adapted representation is consumed by an MPC planner and by embedding-space heads, the adaptation signal is a grounded latent *prediction* loss (a real transition, not a pseudo-label), and the quantity that must not move is the meaning of the latent coordinates themselves. We show that in this closed-loop regime the failure mode is not collapse or forgetting but silent re-mapping — invisible to the adaptation loss, fatal to goal embeddings and frozen heads — and that anchoring the *prediction target* (rather than logits or weights) removes it exactly.

Stop-gradients, EMA teachers, and collapse. BYOL and SimSiam avoid representational collapse with EMA targets and stop-gradients [8, 6]; V-JEPA and TD-MPC2 stabilize latent prediction similarly [4, 12]. LACE transplants this device from *pretraining* (where the teacher and student co-evolve over millions of steps) to *deployment-time adaptation* (where the anchor must guarantee consistency with already-calibrated consumers) — a different requirement that favors a frozen or very slow anchor (Section 5.7).

Adaptive control. Classical adaptive control separates the fixed reference model from the adapting plant model [3]; LACE is the representation-learning analogue: the anchor is the reference frame, the student is the adapting model.

3 Preliminaries: AdaJEPA

Following AdaJEPA authors [1]: an encoder $z_t = E_\theta(o_t)$, an action encoder $u_t = E_a(a_t)$, and a predictor $\hat{z}_{t+1} = f(z_t, u_t)$ are pretrained on reward-free exploration trajectories with the multi-step latent prediction objective plus variance/covariance regularization. At test time the agent plans with MPC (CEM or gradient descent) over the latent rollout cost $\sum_k \alpha_k \|\hat{z}_{t+k} - z_g\|^2$ toward a goal latent $z_g = E_\theta(o_g)$, executes a chunk of actions, and adapts on each observed transition with Equation (1) over a recent-transition buffer, updating only a parameter subset (by default the last predictor block and the encoder projection head, with the encoder learning rate $30\times$ smaller). Weights are episode-local: each episode starts from the pretrained checkpoint.

4 The head-consistency failure, and LACE

4.1 Frozen consumers make the failure visible

The adaptation loss of Equation (1) is computed *inside* the adapting representation, so it cannot detect a coordinated re-mapping of that representation. To measure such re-mapping on public benchmarks we attach *frozen probe heads*: after pretraining, we freeze the world model and train small MLPs on its latents — (i) a success classifier $p(\text{goal reached} \mid z_t, z_g)$, (ii) a progress regressor estimating the negative task distance, (iii) a state readout (agent position; block pose). The probes are then frozen forever. At evaluation, each arm feeds the probes the latents it actually plans with. A TTA method that merely re-expresses the pretrained representation keeps probe metrics at frozen-model levels; one that relocates the space degrades them — exactly the quantity a deployed system cares about, since success detectors, safety monitors, and logging classifiers are all trained offline on the pretrained space.

4.2 LACE: anchor the target, anchor the goal

LACE replaces the stop-gradient student target with an *anchor* encoder $E_{\bar{\theta}}$: Equation (2), where $\bar{\theta}$ is either (a) the pretrained encoder weights, frozen (**laced-frozen**), or (b) a slow EMA $\bar{\theta} \leftarrow \tau\theta + (1-\tau)\theta$ (**laced-ema**). The anchor adds one encoder forward per adaptation step (the target branch already required one) and no planning-time cost.

Two properties follow. *Anchored consistency*: the student can only reduce the loss by predicting where the *pretrained* encoder will place the next observation — modeling dynamics in the fixed reference frame — so latent-space relocation is penalized rather than rewarded. *Goal stability*: since the anchor defines the reference frame, the goal latent should be computed by it, $z_g = E_{\bar{\theta}}(o_g)$. Re-encoding the goal with the adapting student (the AdaJEPA default) makes the planner chase a moving target: we measure the drift $\|z_g^{(t)} - z_g^{(0)}\|$ directly in Section 5.5.

The name is the method: adaptation that has “no pretensions to originate anything.” The student may learn new dynamics at full learning rate, but the meaning of the coordinates is inherited, not invented.

5 Experiments

We evaluate four arms throughout: **frozen** (no adaptation), **unlaced** (AdaJEPA’s self-referential target, Eq. 1), **laced-frozen** (frozen anchor, Eq. 2), and **laced-ema** (EMA anchor, $\tau=0.996$). Unless noted, results use CEM planning, 15 episodes $\times$ 2 seeds per (shift, arm) cell, and the AdaJEPA adaptation recipe (predictor tail + encoder head, asymmetric learning rates).

5.1 Setup: miniature reference benchmarks

PushObj-mini. A 2D pusher (velocity-controlled disc) and a rigid polyomino block on a bounded table (pymunk contact dynamics), rendered to 64×64 RGB with a 2-frame history stack. Train shapes {T, L, Z, +}; held-out test shapes {I, smallT, cube}. Reward-free pretraining data uses contact-biased Ornstein–Uhlenbeck exploration (trajectories re-rolled until the block moves; a 30% pure-OU slice preserves “no contact,




Figure 1: PointMaze high-damping shift, four adaptation recipes $\times$ two target sources. Anchoring the target costs nothing: laced (teal) matches unlaced (red) planning success at every recipe (left), and within short episodes at conservative rates neither damages the frozen success head (right). The damage appears under pressure — Sections 5.4, 5.5, and 5.8.

no motion” examples). Goals are reachable states 8 steps ahead under the data policy; success requires block pose within 0.5 units and 0.6 rad (modulo each shape’s rotational symmetry). Tasks are verified oracle-solvable (8/8 by ground-truth-state CEM), but the learned world model plans at 10–25% success on this benchmark and the pre-registered seen-vs-unseen generalization gate did not pass (gap 2.2 pp, threshold 10 pp); we therefore use PushObj for the drift and anchor-speed analyses (E4, E6), where its longer horizons are informative, and do not rest planning claims on it.

PointMaze-mini. A point mass in procedurally generated mazes with dynamics shifts (mass $\times 0.2$, damping $\times 8$), visual corruptions (blur, salt-and-pepper, darkening, agent recolor), and held-out layouts.

Both use the same small CNN encoder (128-d latent), residual-MLP predictor, CEM/GD planners, and the AdaJEPA adaptation recipe. Pre-registered pass/fail gates for every claim were committed before the grid ran (GATES.md in the code release).

5.2 E1: adaptation gains transfer across target sources

On the high-damping maze shift, adaptation lifts planning success from 56.7% (frozen) to 83.3% with the full recipe, and the lift is *identical* whether the target is the student or the frozen anchor (Figure 1, left): anchored consistency does not tax adaptation. Probe AUC stays near the frozen level (0.96–0.98) for all arms — with the paper’s conservative asymmetric learning rates and short episodes, the maze alone does not expose the failure. Three stressors do: goal-latent drift (E4), symmetric learning rates (E3), and long-horizon deployment (E7).

5.3 E2: full shift suites

Across the maze suite (Figure 2), the laced arms match or beat unlaced planning success in aggregate — dynamics shifts: unlaced 90.0%, laced-frozen 88.3%, frozen 78.3%; visual shifts: laced-frozen 72.5% vs unlaced 71.7%, frozen 66.7%; held-out layouts: laced-ema 44.4% vs unlaced 45.6%, frozen 35.6%. With the gradient-descent planner on dynamics shifts the anchored arm is strongest (90.0% vs 70.0% unlaced, 73.3% frozen): GD planning is more sensitive to the smoothness of the cost landscape, which the anchor preserves. On PushObj-mini the planner operates near its noise floor (Section 5.1; success 10–25% everywhere), and per-arm differences are within error bars; we report the suite for completeness and lean on the maze and deployment data for the planning claims.




Figure 2: PointMaze shift suite (dynamics + visual), all four arms.

5.4 E3: symmetric learning rates and collapse safety

AdaJEPA adapts the encoder $30\times$ slower than the predictor; this asymmetry is the standard defense against self-referential drift. E3 removes it: encoder and predictor share one learning rate, swept over $\{10^{-5}, 10^{-4}, 3\times 10^{-4}, 10^{-3}\}$. The unlaced arm improves until 3×10^{-4} and then collapses at 10^{-3} (33.3% success, worse than never adapting). The laced arm is monotone in the sweep and *best* at the most aggressive rate (86.7%): with the target pinned, a larger step size only means faster convergence to the anchored solution, not faster relocation. Anchoring replaces a fragile hyperparameter with a structural guarantee.

5.5 E4: goal-latent drift

The cleanest signature of relocation needs no probes: re-encode the *same* goal observation at every replanning step and measure how far its latent moves. Unlaced drifts steadily (mean final drift 0.11 on PushObj, 0.21 on the maze suite — roughly 8% of the typical inter-state latent distance); laced-frozen is exactly zero by construction, and laced-ema stays two orders of magnitude smaller (0.008). The planner in the unlaced arm is chasing a target that moves because the map, not the territory, is changing.

5.6 E5: low-data pretraining

Sweeping pretraining shape diversity $K \in \{1, 2, 4\}$ on PushObj and evaluating on unseen shapes, adaptation helps most exactly where data is scarcest (unlaced +laced both lift $K=1, 2$ success), consistent with AdaJEPA’s data-compensation result; anchored and self-referential targets are again indistinguishable in success. Probe AUC in this regime is dominated by the quality of the pretrained representation itself (frozen-arm AUC falls from 0.96 at $K=4$ to 0.67 at $K=1$), so head-consistency comparisons at $K=1$ mostly measure the probes’ failure to fit, not adaptation damage; we flag this as a limit of the probe methodology at very low data.

5.7 E6: how slow must the anchor be?

Sweeping the EMA decay $\tau \in \{0.9, 0.996, 1.0\}$ on unseen PushObj shapes: a fast anchor ($\tau=0.9$) is self-reference in slow motion — probe AUC falls to 0.68 and planning success to 5% — while $\tau=0.996$ and the frozen anchor hold AUC at 0.85–0.91. Consistency requires the anchor to move at most negligibly over the horizon on which downstream consumers were calibrated; when in doubt, freeze it.




Figure 3: Symmetric-LR sweep on the high-damping shift. At $\eta_{\text{enc}}=10^{-3}$ the self-referential target collapses planning (33.3%, below the frozen baseline) while the anchored target keeps improving (86.7%).

5.8 E7: deployed screen-interaction agent

The deployment replay is where the phenomenon was first observed, and it is also where the simple story breaks — we report both halves. On 110 recorded runs of a screen-interaction agent (5,434 transitions, prequential scoring: each step is evaluated before the model adapts on it), the published full recipe cuts second-half prediction divergence by 44.6% while the deployed success head’s AUC on predicted latents falls from 0.933 to 0.826. Anchoring the target recovers only a fraction (0.838) — our pre-registered gate for this experiment (AUC within 0.02 of frozen) *fails*.

The parameter-subset decomposition (Figure 6) explains why. Head damage here has two mechanisms: (i) *encoder relocation*, the self-reference failure this paper is about, which anchoring removes completely — adapting only the encoder with the self-referential target drops pred-latent AUC to 0.914 and encoder-space AUC below baseline (0.674 vs 0.695), while the same adaptation with an anchored target restores both exactly to frozen levels (0.933, 0.691); and (ii) *predictor drift* — adapting only the predictor drops AUC to 0.840 *identically* under both target sources, because the anchor constrains the target, not the predictor’s overfitting to the small recent-transition buffer. LACE is therefore necessary but not sufficient at deployment scale: it solves the representation half of the problem exactly, and sharply exposes the optimization half (buffer size, step count, trust regions on the predictor), which we leave open.

6 Discussion and limitations

The anchor-drift tension. A frozen anchor guarantees consistency but can never incorporate a truly novel world; under persistent shift the anchored target is eventually wrong about *content*, not just coordinates. The EMA anchor trades consistency for plasticity (E6); scheduling that trade — and re-calibrating consumers when the anchor must move — is the continual learning question this paper deliberately does not solve.

Scale. Our benchmarks are miniature by design (minutes per training run on a laptop GPU); the deployment replay is the evidence that the phenomenon and the fix survive contact with reality. Scaling LACE to foundation-scale world models is future work.

When unlaced TTA is fine. If nothing downstream is frozen — no fixed goal embeddings, no calibrated heads, single-episode tasks scored only by environment reward — self-referential TTA’s extra freedom can be harmless at conservative learning rates (our maze grids show small recipes cause little damage within short episodes). The failure grows with adaptation pressure: step count, learning rate, and encoder participation.




Figure 4: Goal-latent drift $\|z_g^{(t)} - z_g^{(0)}\|$ pooled over the PushObj suite. The goal observation never changes; any drift is encoder relocation.

Reproducibility statement

All environments, training, adaptation, evaluation, probes, and figures are released with exact seeds and one-command reproduction scripts. The deployment replay data is proprietary telemetry and is not released; we report aggregate metrics only, and the public benchmarks reproduce the phenomenon end-to-end.

Acknowledgments

We thank the AdaJEPA authors for a reproducible and clearly written method, and Ada Lovelace for the epigraph and the acronym.¹

References

- [1] AdaJEPA authors. AdaJEPA: Test-time adaptation of joint-embedding predictive world models for planning under distribution shift. *arXiv preprint arXiv:2606.32026*, 2026. Reference benchmark and baseline for this paper.
- [2] Mahmoud Assran, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, Yann LeCun, and Nicolas Ballas. Self-supervised learning from images with a joint-embedding predictive architecture. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023.
- [3] Karl J Åström and Björn Wittenmark. *Adaptive Control*. Dover, 2nd edition, 2008.
- [4] Adrien Bardes, Jean Ponce, Yann LeCun, et al. V-JEPA: Latent video prediction for visual representation learning. *arXiv preprint*, 2024.
- [5] Goirik Chakrabarty, Manogna Sreenivas, and Soma Biswas. SATA: Source anchoring and target alignment network for continual test time adaptation. *arXiv preprint arXiv:2304.10113*, 2023.
- [6] Xinlei Chen and Kaiming He. Exploring simple siamese representation learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2021.

¹LACE: Learning Anchored Consistent Embeddings. The model is LovelaceJEPA because it is an Ada-JEPA that refuses to originate anything.




Figure 5: Prequential replay of 110 deployed runs (5,434 steps): each transition is scored before being adapted on. Divergence gains (left) survive anchoring; deployed-head AUC (right) recovers only partially, motivating the mechanism decomposition (Figure 6).

- [7] Yossi Gandelsman, Yu Sun, Xinlei Chen, and Alexei A Efros. Test-time training with masked autoencoders. In *Advances in Neural Information Processing Systems*, 2022.
- [8] Jean-Bastien Grill, Florian Strub, Florent Altché, Corentin Tallec, Pierre Richemond, Elena Buchatskaya, Carl Doersch, Bernardo Avila Pires, Zhaohan Guo, Mohammad Gheshlaghi Azar, et al. Bootstrap your own latent: A new approach to self-supervised learning. In *Advances in Neural Information Processing Systems*, 2020.
- [9] David Ha and Jürgen Schmidhuber. World models. *arXiv preprint arXiv:1803.10122*, 2018.
- [10] Danijar Hafner, Timothy Lillicrap, Jimmy Ba, and Mohammad Norouzi. Dream to control: Learning behaviors by latent imagination. In *International Conference on Learning Representations*, 2020.
- [11] Nicklas Hansen, Xiaolong Wang, and Hao Su. Temporal difference learning for model predictive control. In *International Conference on Machine Learning*, 2022.
- [12] Nicklas Hansen, Hao Su, and Xiaolong Wang. TD-MPC2: Scalable, robust world models for continuous control. In *International Conference on Learning Representations*, 2024.
- [13] Yann LeCun. A path towards autonomous machine intelligence. *OpenReview preprint*, 2022.
- [14] Shuaicheng Niu, Jiaxiang Wu, Yifan Zhang, Yafo Chen, Shijian Zheng, Peilin Zhao, and Mingkui Tan. Efficient test-time model adaptation without forgetting. In *International Conference on Machine Learning*, 2022.
- [15] Junha Song, Jungsoo Lee, In So Kweon, and Sungha Choi. EcoTTA: Memory-efficient continual test-time adaptation via self-distilled regularization. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023.
- [16] Yu Sun, Xiaolong Wang, Zhuang Liu, John Miller, Alexei Efros, and Moritz Hardt. Test-time training with self-supervision for generalization under distribution shifts. In *International Conference on Machine Learning*, 2020.
- [17] Dequan Wang, Evan Shelhamer, Shaoteng Liu, Bruno Olshausen, and Trevor Darrell. Tent: Fully test-time adaptation by entropy minimization. In *International Conference on Learning Representations*, 2021.




Figure 6: Parameter-subset ablation on the replay. Adapting only the predictor already reproduces most of the pred-latent AUC damage, and anchoring cannot prevent it (the anchor constrains the *target*, not predictor overfitting to the small transition buffer). Encoder-space AUC (right) isolates the relocation component that anchoring does remove.

- [18] Qin Wang, Olga Fink, Luc Van Gool, and Dengxin Dai. Continual test-time domain adaptation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022.
- [19] Hao Zhao, Yuejiang Liu, Alexandre Alahi, and Tao Lin. On pitfalls of test-time adaptation. *International Conference on Machine Learning*, 2023.
- [20] Gaoyue Zhou, Hengkai Pan, Yann LeCun, and Lerrel Pinto. DINO-WM: World models on pre-trained visual features enable zero-shot planning. In *arXiv preprint arXiv:2411.04983*, 2024.

A Environment details

B Probe training details

C Full result tables